CurveAttentionBlock: the exact pre-LN update

Repeated 8 times. Only Q changes between blocks; the same M is supplied to every cross-attention layer. Shapes omit the pair batch axis P.

One block

Every attention module has 8 heads, model width 512, and attention dropout 0.1. The circles are residual additions.

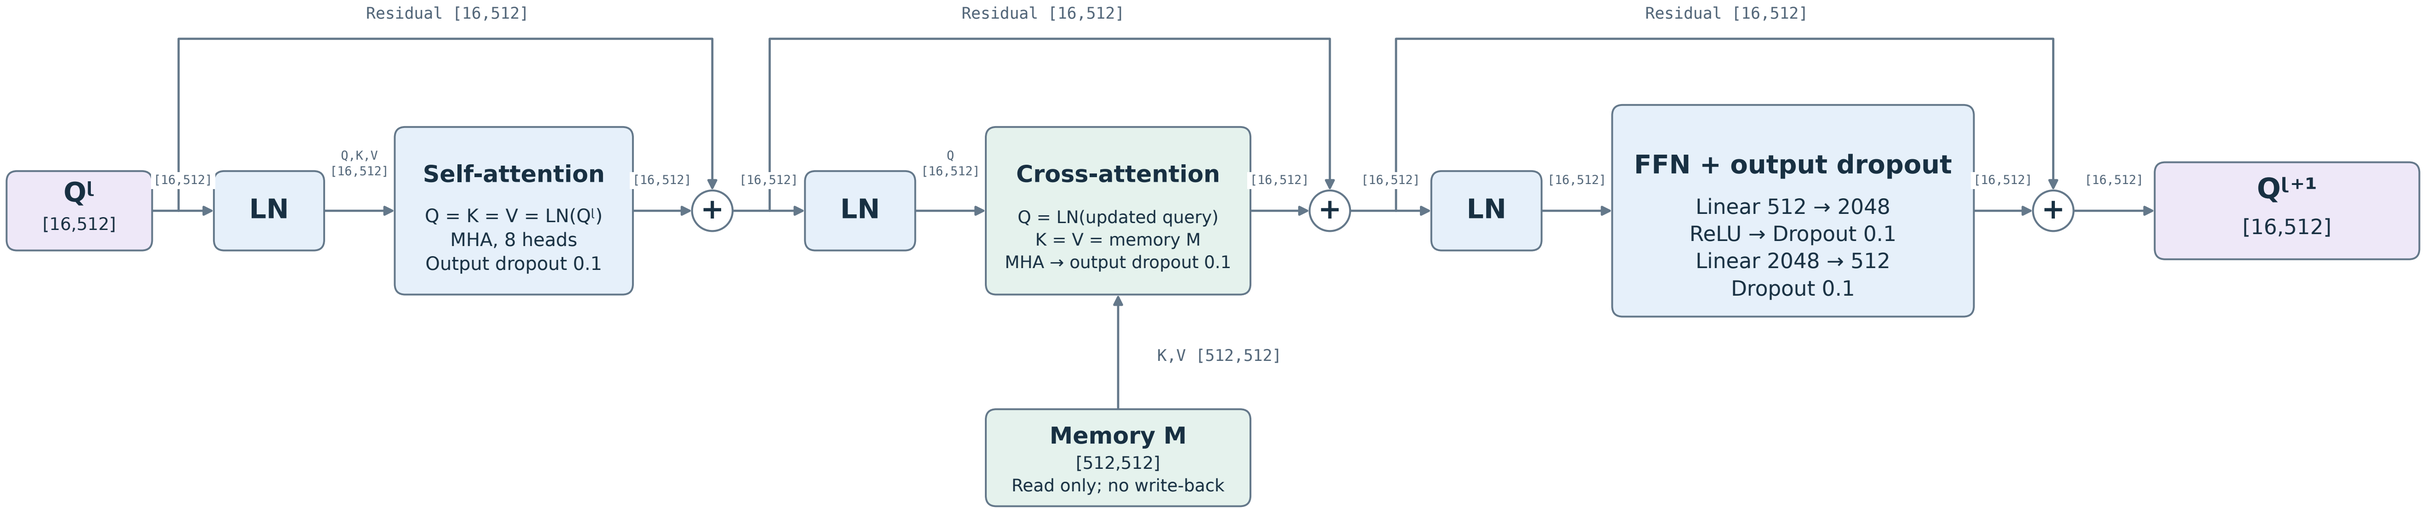

The final block output is $E = Q^8$. DualUVHead consumes E directly; no face feature or memory concatenation is inserted before that head.